

A woman with dark hair, wearing a blue patterned scarf and a brown jacket, is looking down at a smartphone. The background is a blurred mountain landscape.

arm

Chihuahua: A WASM sandbox

Derek Miller, Dominic Mulligan, Shale Xiong, Hugo Vincent
Arm Research

Berkeley WASM outside the browser workshop
9 August 2019

Chihuahua motivation

Part of a wider project (“*Veracruz*”) looking at secure processing environments (SPEs)

Goal: sandbox for running small loadable programs inside SPEs

- Some issues with measurement/executable pages on some SPEs
- Provide confidentiality guarantees for loaded program
- Protection of the host machine from loaded program (e.g. Rowhammer)
- Ability to mitigate architectural vulnerabilities without forcing program recompilation
- Enforce dynamic policies on code running in sandbox
- High-level of trust in the sandbox

Some security engineering principles

Adopt some design principles when developing Veracruz:

- All code should be easily auditable by humans
- Every line of code in the SPE should have a purpose
- Any code not strictly required to be in the SPE should be moved outside

Most code running in SPE is written in Rust. Exception: Chihuahua which is written in C...

- Static analysis used to obtain memory safety, rather than property of language
- Also want to verify properties of Chihuahua

C is (perversely!) one of the best languages to use for verification: it's unsafe, everyone knows it's unsafe, and it's everywhere: so lots of tools exist

Chihuahua, concretely

Chihuahua is a “hardware-like” VM (in JVM sense) implementing Zocalo ISA:

- Half-way house between WASM stack machine and RISC load/store architecture
- Fixed-width instruction encoding (32-bits), load-store architecture
- Separate integer, floating point GPRs in both 32- and 64-bit variants

Most opcodes and semantics borrowed from WASM:

- WASM has structured control-flow, Zocalo de-structures this using jumps and condition flags
- Interpreter main loop simpler to audit and verify (just fetch/execute, no control flow)

Zocalo Virtual Machine:

- Harvard Architecture: separate code, main memories, only fetch reads from code memory
- Parameters passed via operand stack, only push and pop manipulate this
- Separate return address stack: only `call` and `return` manipulate this
- Deterministic: no undefined behaviours, meaning anything strange causes immediate program termination

Why this approach?

WASM is designed for JITing: this is how most (all?) browsers use WASM

- JITs hard to audit, especially sophisticated ones: as line-count increases probability of memory errors, security holes increases. Chihuahua is simple...
- Some SPEs don't sit neatly with JITs
- Dynamic policies easier to enforce with Chihuahua, e.g. resource counts, trace properties of program
- Potentially want to verify WASM to Zocalo translation making use of e.g. Conrad Watt's work in Isabelle/HOL

But...

- We leave potentially a lot of performance on the table
- One idea: we could try writing a high-assurance WASM JIT (interest in collaborating?)

Why WASM in the middle?

Many compilers targeting WASM: use it as an IR and a path from C to Zocalo

Translation tool, `wasm2zocalo`, reuses reference interpreter written in OCaml

- Reuses front-end, parses WASM binary and produces OCaml AST of the WASM program
- Walk the AST translating into Zocalo instructions
- Separate assembler that does some peephole optimisations, generates addresses etc
- Translated 85% of WASM instruction set

Things to be done:

- Control-flow destructuring,
- Setup of subprocedure environment prior to function call,
- Hoist large immediate constant loads into memory loads,
- Add explicit pushes and pops
- Convert 32-bit addresses to 64-bit addresses

Zocalo implementation

Zocalo is implemented in pure C, around 6,000 lines of heavily-commented code

- Merge requests code-reviewed by another team member,
- Compiled with `clang` with all warnings and errors turned on

Use **C Bounded Model Checker (CBMC)** on Zocalo implementation:

- Capable of verifying properties of C code (pre/post-conditions, asserts etc)
- Also detecting undefined behaviours, code-smells in C (e.g. overflow, memory/indexing errors)

Also run Frama-C, a static-analysis/verification platform on code

- Disadvantage: requires extensive EACSL annotations embedded in code
- Mostly does what CBMC does, but more complex to use and understand, so capability a little bitrotted

Other potential avenues: TIS interpreter and Cerberus C tools from Peter Sewell's group

- Arm is a TIS licensee...

What properties do we verify?

Around 10,000 lines of C code for CBMC verification testbenches:

- Verify that encode and decode are correct (mutual partial inverses)
- Verify Zocalo semantics is as expected [†] Floating-point instructions cause CBMC to die ☹️
- Verify machine invariants are preserved by execution: e.g. PC always valid
- Memory, stack read and write functions have correct properties
- Unreachable code paths are marked, unreachability checked for valid inputs
- All Zocalo instructions in execute, decode functions are handled using cover properties

Takes 72 hours to run full test bench:

- 90% of time spent verifying execution
- Flushed out lots of bugs


arm

Translating WASM bytecode to Zocalo

WASM to Zocalo overview

```
(func (;0;) (type 0) (param i32 i64) (result i64)
  (local i32 i32 i64 i64)
  local.get 1
  i64.const 0xffffffff
  i64.ne
  if ;; label = @1
    loop ;; label = @2
      .....
    loop ;; label = @3
      local.get 2
      .....
      i64.lt_u
      br_if 1 (;@2;)
    end
  call 1
  .....
end
end
local.get 4)
```

(func(0,1) 18.Arm.) limited



WASM to Zocalo

```
(func (;0;) (type 0) (param i32 i64) (result i64)
```

```
  (local i32 i32 i64 i64)
```

```
  #100 local.get 1
```

```
  #101 i64.const 0xffffffff
```

```
  #102 i64.ne
```

```
  #103 if ;; label = @1
```

```
  #104 loop ;; label = @2
```

```
    .....
```

```
    #132 loop ;; label = @3
```

```
      local.get 2
```

```
      .....
```

```
      i64.lt_u
```

```
      br_if 1 (;@2;)
```

```
    end
```

```
    call 1
```

```
    .....
```

```
  end
```

```
end
```

```
  local.get 4)
```

```
(func (;1;) .....)
```



Label Table

Label_0 : #43

...

Label_i : #n

...

Label_22 : #103

Label_23 : #104

Label_24 : #132

Label Stack

Label_24

Label_23

Label_22

WASM to Zocalo

```
(func (;0;) (type 0) (param i32 i64) (result i64)
  (local i32 i32 i64 i64)
  #100 local.get 1
  #101 i64.const 0xffffffff
  #102 i64.ne
  #103 if ;; label = @1
  #104 loop ;; label = @2
    .....
  #132 loop ;; label = @3
    local.get 2
    .....
    i64.lt_u
    br_if 1 (;@2;)
  end
  call 1
  .....
end
end
local.get 4)
```



Label Table

Label_0 : #43

...

Label_i : #n

...

Label_22 : #103

Label_23 : #104

Label_24 : #132

Label Stack

Label_24

Label_23

Label_22

Pseudo Zocalo instructions

br_if label_23

WASM to Zocalo

(func (;0;) (type 0) (param i32 i64) (result i64))

R31 points to parameters and local variables stack. Initially it pointed to end of heap.

```
#100 local.get 1
#101 i64.const 0xffffffff
#102 i64.ne
#103 if ;; label = @1
#104 loop ;; label = @2
.....
#132 loop ;; label = @3
  local.get 2
  .....
  i64.lt_u
  br_if 1 (;@2;)
end
call 1
.....
end
end
local.get 4)
```

Label Table

Label_0 : #43

...

Label_i : #n

...

Label_12 : #323

...

Label_k : #m

Pseudo Zocalo instructions

```
// reserve space for local variables
sub r31 r31 0x1c
// pop and store parameters
pop r0
sub r31 r31 0x4
store r0 [r31+0x0]
...
// call the function
call label_12
```



The Arm trademarks featured in this presentation are registered trademarks or trademarks of Arm Limited (or its subsidiaries) in the US and/or elsewhere. All rights reserved. All other marks featured may be trademarks of their respective owners.

www.arm.com/company/policies/trademarks